

STT-TTS and Internet

Multimedia Networking - AI

P. Bakowski



bako@ieee.org



AI - Multimedia networking

prompt :

Write GStreamer pipeline to send **uncompressed** audio over RTP and UDP protocols.

```
gst-launch-1.0 -v autoaudiosrc ! audioconvert !  
audio/x-raw,format=S8,channels=1,rate=16000 !  
rtpL8pay ! udpsink host=<receiver_ip> port=<port>
```

compression ratio = 1

prompt :

Write GStreamer pipeline to send **compressed MP3** audio over RTP and UDP protocols.

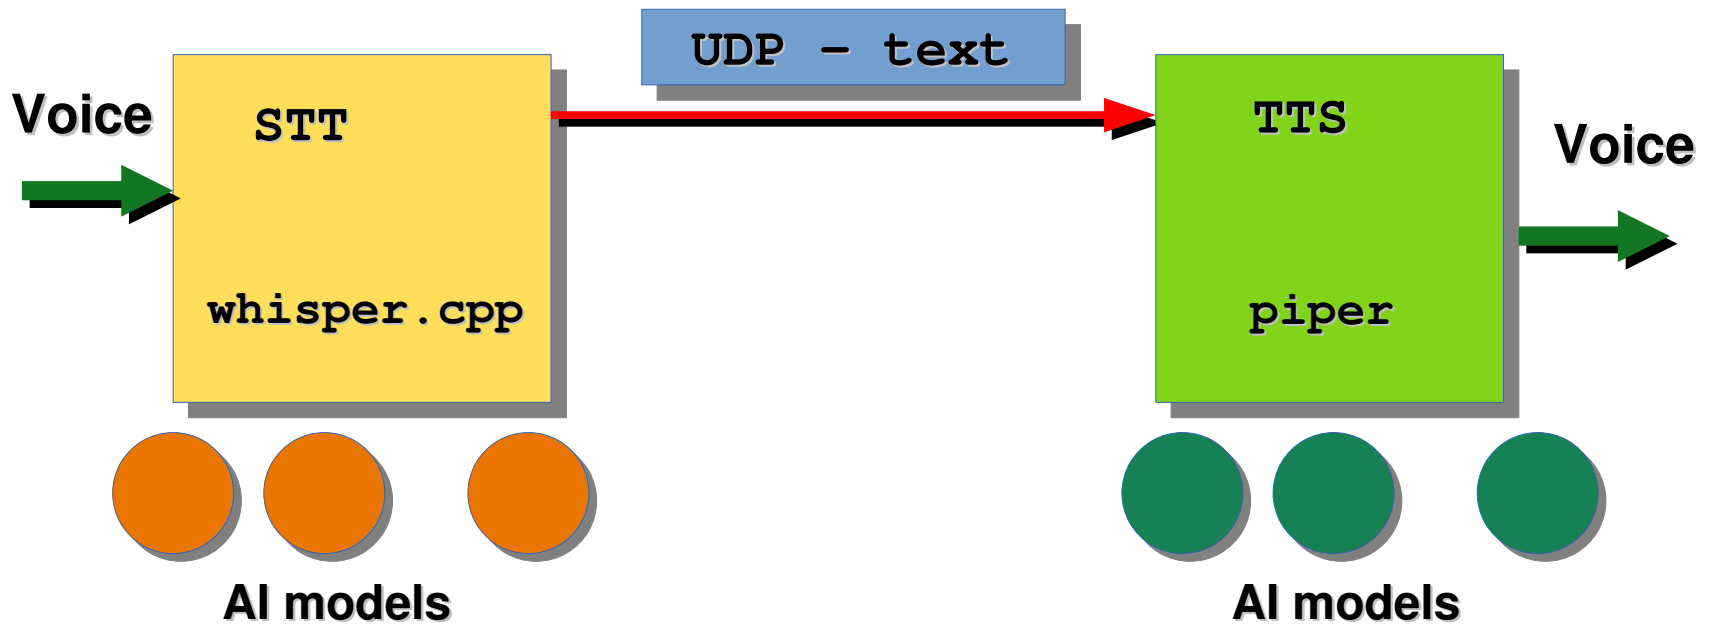
```
gst-launch-1.0 -v autoaudiosrc ! audioconvert !  
lamemp3enc ! rtpmpapay ! udpsink host=<receiver_ip>  
port=<port>
```

compression ratio ~= 5

AI - Multimedia networking

ASP – Automatic Speech Recognition – **Speech To Text**

Text To Speech – Text to Voice



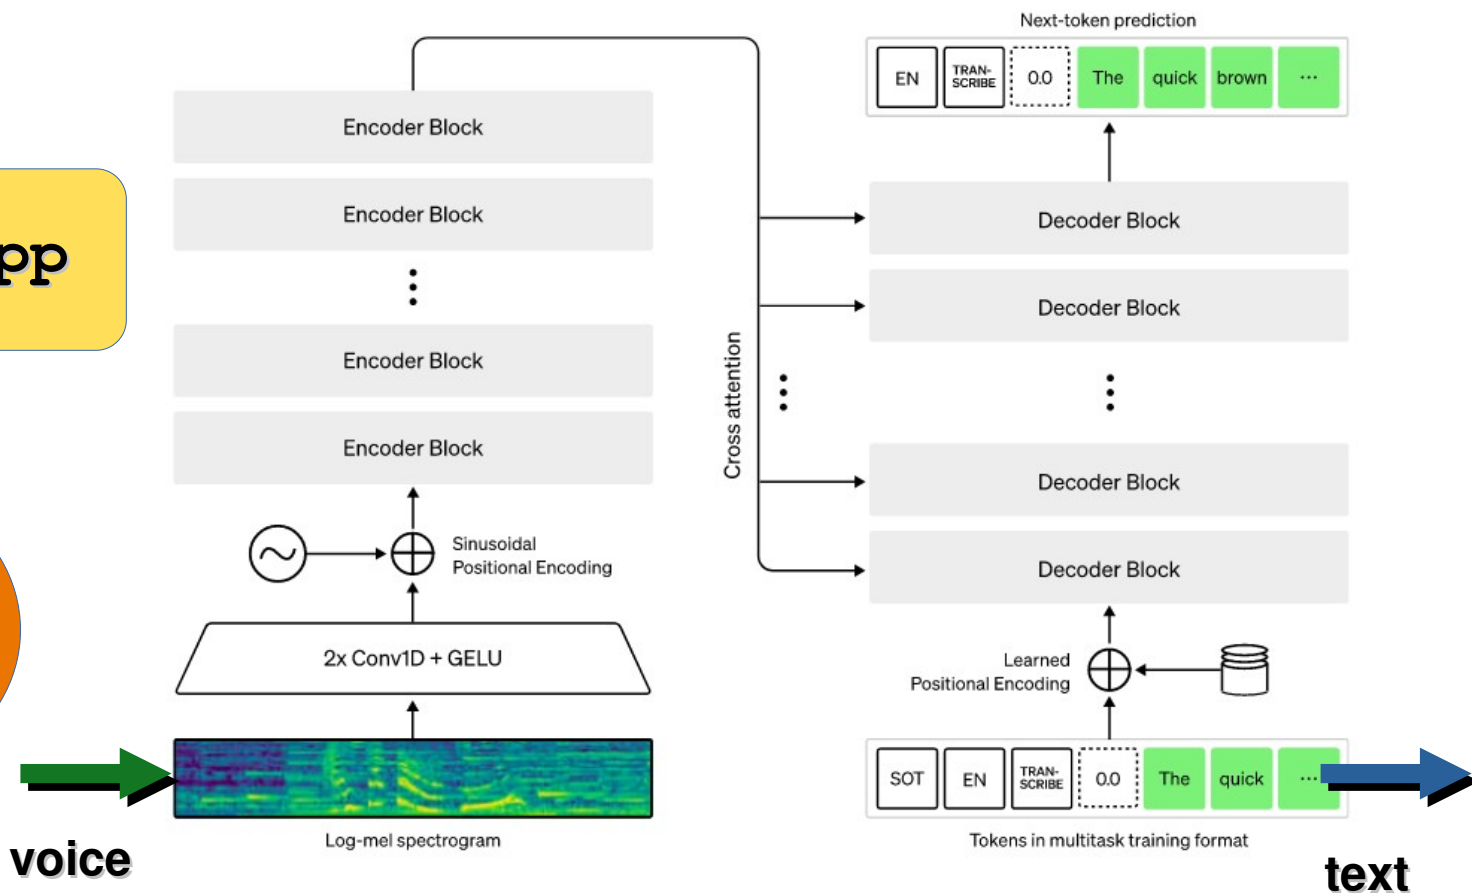
“compression” ratio ~ 50

AI – Speech To Text : `whisper.cpp`

The **whisper architecture** is a simple end-to-end approach, implemented as an **encoder-decoder**

`whisper.cpp`

model



AI – Speech To Text : whisper.cpp

whisper.cpp is the framework – ./stream is executable

```
rock@rock-5a:~/whisper.cpp$ ./stream -h
```

```
usage: ./stream [options]
```

```
options:
```

-h,	--help	[default]	show this help message and exit
-t N,	--threads N	[4]	number of threads to use during computation
	--step N	[3000]	audio step size in milliseconds
	--length N	[10000]	audio length in milliseconds
	--keep N	[200]	audio to keep from previous step in ms
-c ID,	--capture ID	[-1]	capture device ID
-mt N,	--max-tokens N	[32]	maximum number of tokens per audio chunk
-ac N,	--audio-ctx N	[0]	audio context size (0 - all)
-vth N,	--vad-thold N	[0.60]	voice activity detection threshold
-fth N,	--freq-thold N	[100.00]	high-pass frequency cutoff
-su,	--speed-up	[false]	speed up audio by x2 (reduced accuracy)
-tr,	--translate	[false]	translate from source language to english
-nf,	--no-fallback	[false]	do not use temperature fallback while decoding
-ps,	--print-special	[false]	print special tokens
-kc,	--keep-context	[false]	keep context between audio chunks
-l LANG,	--language LANG	[en]	spoken language
-m FNAME,	--model FNAME	[models/ggml-base.en.bin]	model path
-f FNAME,	--file FNAME	[text output file name
-tdrz,	--tinydiarize	[false]	enable tinydiarize (requires a tdrz model)
-sa,	--save-audio	[false]	save the recorded audio to a file
-ng,	--no-gpu	[false]	disable GPU inference

AI – Speech To Text : whisper.cpp

```
./stream -m models/ggml-base.en.bin -ac 1024
```

```
./stream -m models/ggml-base.en.bin --step 2000 --length 8000 \  
-c 0 -t 6 -ac 512 -vth 0.6
```

```
rock@rock-5b:~/rockAI/whisper.cpp$ ./stream -m models/ggml-base.en.bin -ac  
1024
```

```
init: found 3 capture devices:
```

```
init:   - Capture device #0: 'Built-in Audio Stereo'
```

```
init:   - Capture device #1: 'HD Pro Webcam C920 Analog Stereo'
```

```
init:   - Capture device #2: 'Built-in Audio Stereo (2)'
```

```
init: attempt to open default capture device ...
```

```
init: obtained spec for input device (SDL Id = 2):
```

```
init:   - sample rate:      16000
```


```
init:   - format:          33056 (required: 33056)
```

```
init:   - channels:         1 (required: 1)
```

```
init:   - samples per frame: 1024
```

```
..
```

./stream



AI – Speech To Text : whisper.cpp

```
rock@rock-5b:~/rockAI/whisper.cpp$ ./stream -m models/ggml-base.en.bin -ac 1024
```

```
init: found 3 capture devices:
```

```
init:   - Capture device #0: 'Built-in Audio Stereo'
```

```
init:   - Capture device #1: 'HD Pro Webcam C920 Analog Stereo'
```

```
init:   - Capture device #2: 'Built-in Audio Stereo (2)'
```

```
init: attempt to open default capture device ...
```

```
init: obtained spec for input device (SDL Id = 2):
```

```
init:   - sample rate:      16000
```

```
init:   - format:          33056 (required: 33056)
```

```
init:   - channels:        1 (required: 1)
```

```
init:   - samples per frame: 1024
```

```
..
```

```
[Start speaking]
```


```
[ Silence ]
```

```
(crickets chirping)
```

```
Hello, how are you going?
```

```
[ Silence ]
```


```
Where is the next bus station?
```



./stream – filtering pure speech

```
#include <stdio.h>
#include <string.h>
```

```
int main() {
    int c;
    int txt=0,nl=0,sp=0,sb=0,nb=0,count=0;
    char buff[512];int i=0;
    while ((c = getchar()) != EOF) {
        putchar(c);    // to be commented
        count++;
        if(c==0x0D) { count=0; }
        if(count==2)
        {
            switch(c)
            {
                case 0x20: txt=0;sp=1; break;
                case 0x5B: txt=0;sb=1; break;
                case 0x28: txt=0;nb=1; break;
                default: txt=1;sp=0;sb=0;nb=0;i=0;memset(buff,0x00,512);break;
            }
        }
        if(txt) buff[i++]=c;
        if(c=='.' || c=='?' || c=='!') { txt=0; printf("\n%s\n",buff);}
    }
    return 0;
}
```



filtered "pure
speech" text



. / stream – filtering pure speech

```
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#include <arpa/inet.h>
#include <sys/socket.h>
#define BUFLen 512          // Max length of buffer
#define PORT 8888           // The port on which to send or listen for data
#define REMOTE_ADDR "192.168.1.31"

int main() {
    int c;
    int txt=0,nl=0,sp=0,sb=0,nb=0,count=0;
    char buff[512];int i=0;

    struct sockaddr_in si_me, si_other;
    int s, slen = sizeof(si_other) , recv_len;
    if ((s=socket(AF_INET,SOCK_DGRAM,IPPROTO_UDP))==-1){exit(1);}
    memset((char *) &si_other, 0, sizeof(si_other));
    si_other.sin_family = AF_INET;
    si_other.sin_port = htons(PORT);      // to work remotely
    si_other.sin_addr.s_addr = inet_addr(REMOTE_ADDR);    // to work remotely
```

preparing UDP socket



./stream – filtering pure speech

```
// Read characters from standard input until EOF
```

```
memset(buff,0x00,512);
```

```
while ((c = getchar()) != EOF)
```

```
{
```

```
    putchar(c);
```

```
    count++;
```

```
    if(c==0x0D) { count=0; }
```

```
    if(count==2)
```

```
    {
```

```
        switch(c)
```

```
        {
```

```
            case 0x20: txt=0;sp=1; break;
```

```
            case 0x5B: txt=0;sb=1; break;
```

```
            case 0x28: txt=0;nb=1; break;
```

```
            default: txt=1;sp=0;sb=0;nb=0;i=0;memset(buff,0x00,512);break;
```

```
        }
```

```
    }
```

```
    if(txt) buff[i++]=c;
```

```
    if(c=='.' || c=='?' || c=='!')
```

```
    { txt=0; printf("\n%s\n",buff);
```


```
        sendto(s,buff,strlen(buff)+1,0,(struct sockaddr *)&si_other,slen);
```

```
    }
```

```
}
```


```
return 0;
```

```
}
```



sending UDP packet

./stream .. | sendsock



On the receiving side :

```
..  
while(1)  
{  
    ./stream -m models/ggml-base.en.bin -ac 1024 | ./sendsock  
  
    if((rlen=recvfrom(s,buf,BUFLEN,0,(struct sockaddr *)&si_other,&slen))== -1)  
        {printf("recvfrom error\n"); }  
    printf("%s\n", buf);  
}
```



`echo 'text' | ./piper (TTS)`



`piper` is optimized to perform well on **ARM** (NEON); it takes 1 second to generate 1.6 seconds of speech audio at a medium quality level.

`piper` can stream raw audio to `stdout` as its produced.

`piper` needs pre-trained models (`.onnx`) of the speaker

Streaming audio with `echo` command:

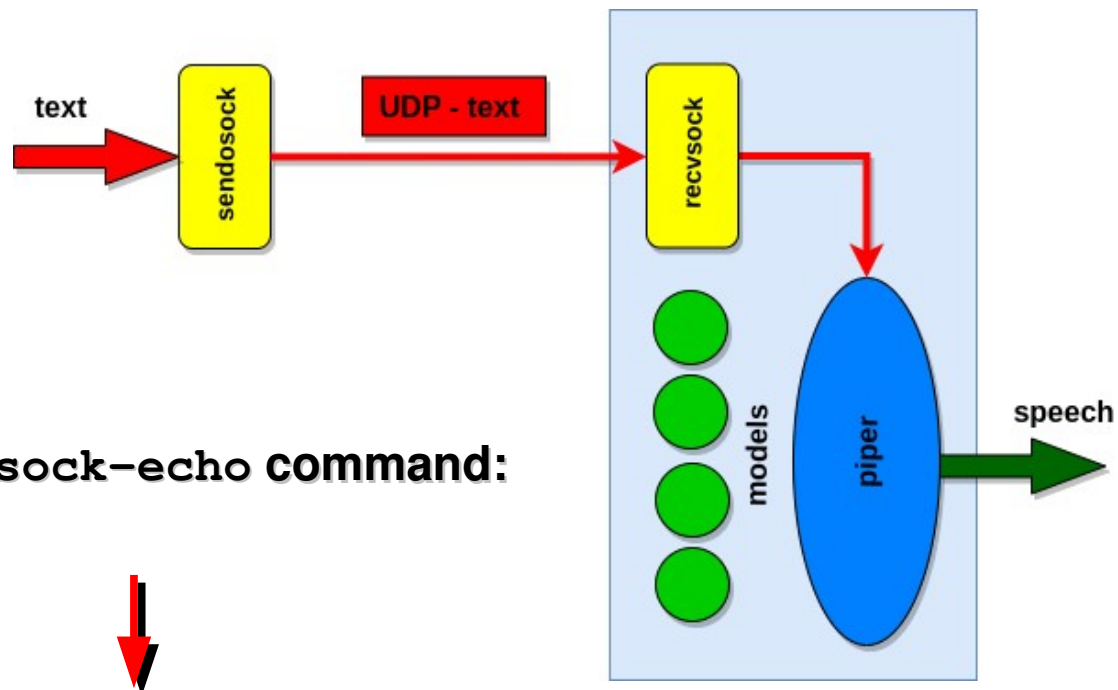
```
echo 'This sentence is spoken first. This is Smart Computer Lab  
multimedia,internet and AI module.' | \  
./piper --model GB_female_south/en_GB-southern_english_female-low.onnx --  
output-raw | \  
aplay -r 15000 -f S16_LE -t raw -
```



`aplay` to reproduce the samples of the voice on audio output

`./recvsock | ./piper`

`./recvsock` is a program that receives UDP datagrams with text and redirects them to the `./piper` via echo command.



Streaming audio with `./recvsock-echo` command:

**`./recvsock | \`
**`./piper --model GB_female_south/en_GB-southern_english_female-low.onnx \`
`--output-raw | aplay -r 15000 -f S16_LE -t raw -`****


./recvsock | ./piper

```
while(1)
{
    memset(message, 0x00, 512);
    if((rlen=recvfrom(s, message, 512, 0, (struct sockaddr *)&si_other, &slen))==-1)
        {printf("recvfrom() error\n"); }
    memset(buffer, 0x00, 1000);
    strcpy(buffer, "echo ");
    strcat(buffer, message);
    sleep(1);
    // generate echo command with message
    retcode=system(buffer);
    if (retcode == 0) { printf("executed OK\n"); }
    else { printf("execution error\n"); }
}
```

echo

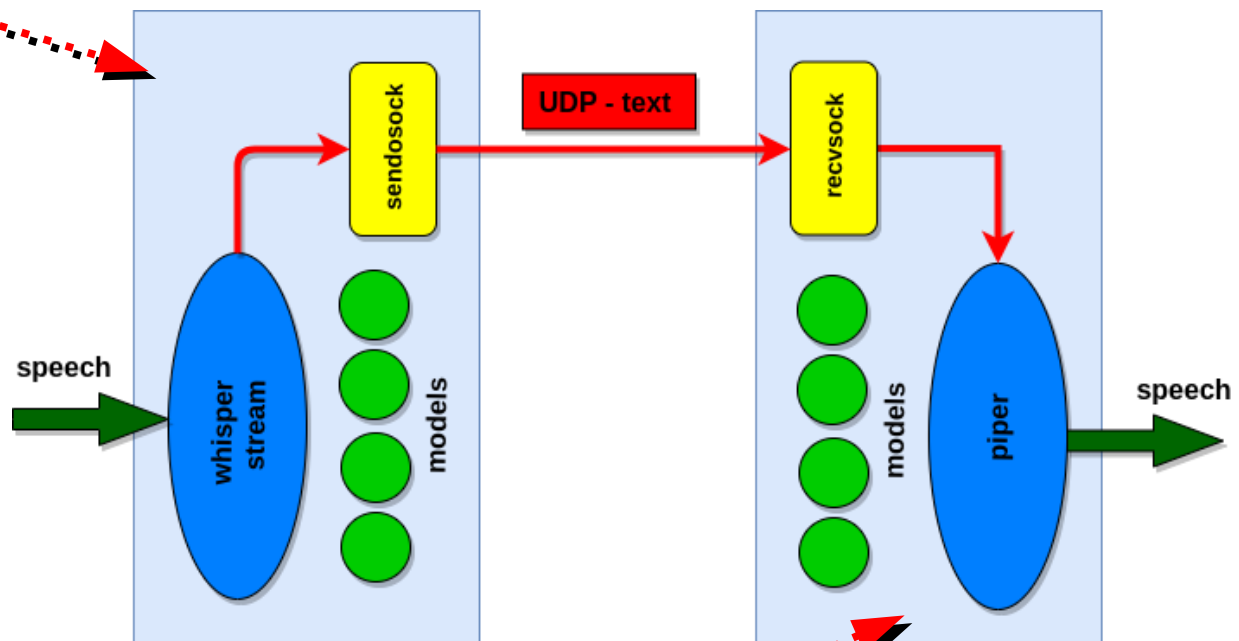


```
./recvsock | \
./piper --model GB_female_south/en_GB-southern_english_female-low.onnx \
--output-raw | aplay -r 15000 -f S16_LE -t raw -
```




whisper.cpp - UDP - piper

```
./stream -m models/ggml-base.en.bin -ac 1024 | ./sendsock
```



```
./recvsock | \  
./piper --model GB_female_south/en_GB-southern_english_female-low.onnx \  
--output-raw | aplay -r 15000 -f S16_LE -t raw -
```

Summary



ASP – Automatic Speech Recognition – **Speech To Text** with `whisper.cpp (./stream)`

Filtering speech from generated text (sounds, noise, silence, and speech)

Sending and receiving text
speech over UDP

Text To Speech with `./piper`

receiving text speech over UDP and playout